1. Kinesis Firehose Delivery Stream

In the AWS Console -> Kinesis

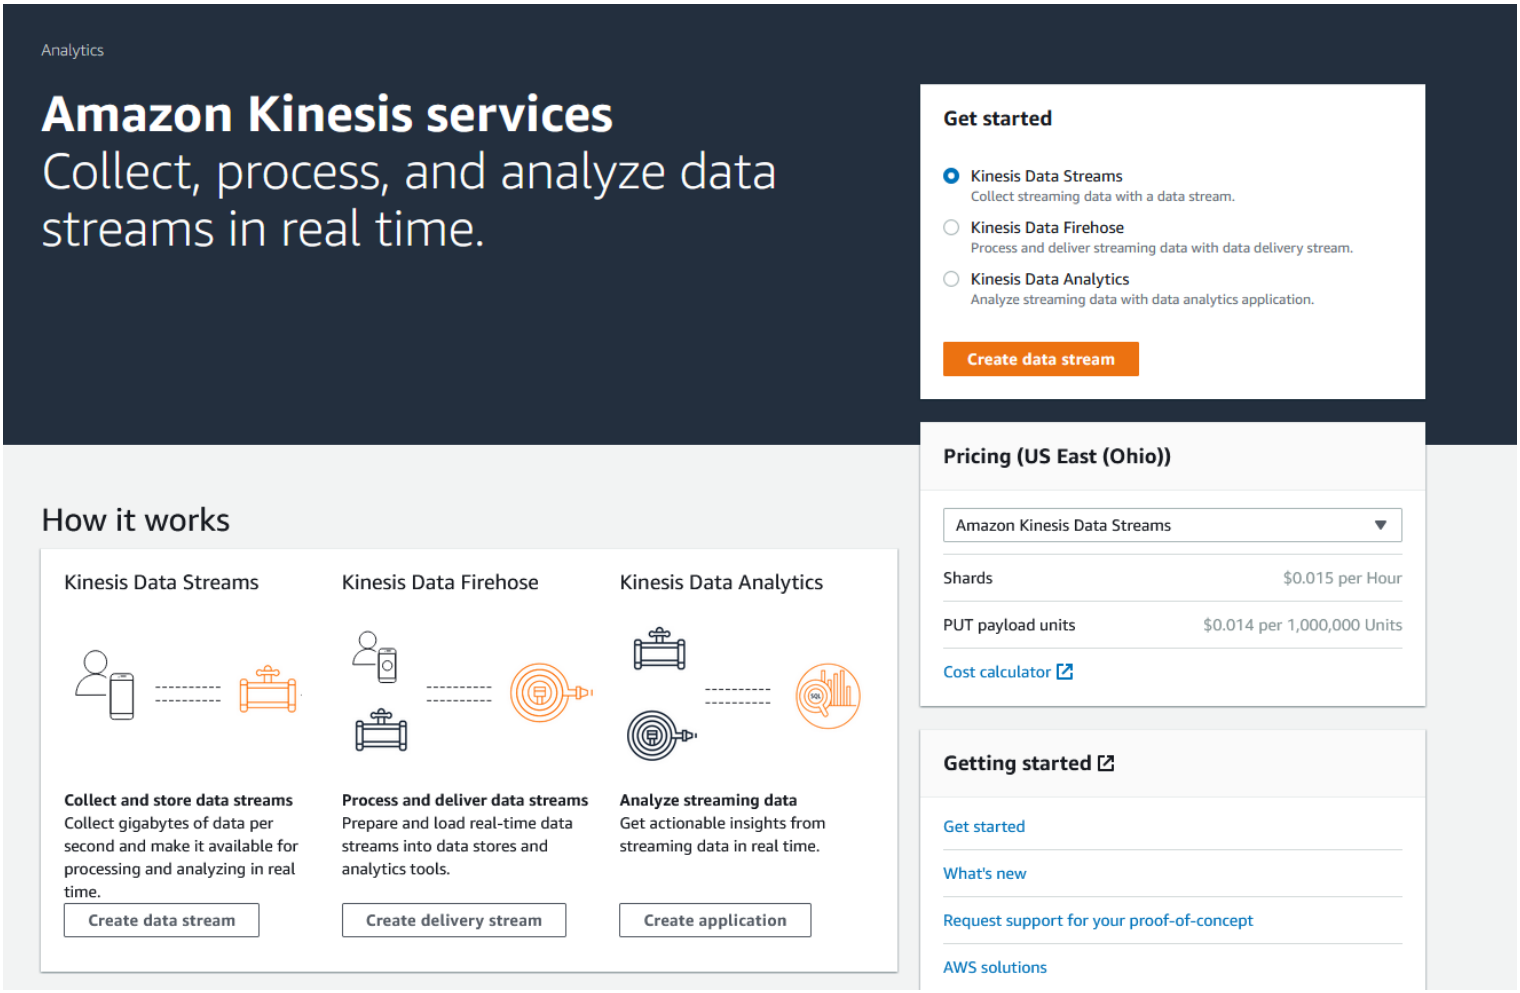

Click “Create Delivery Stream”

Select source as –

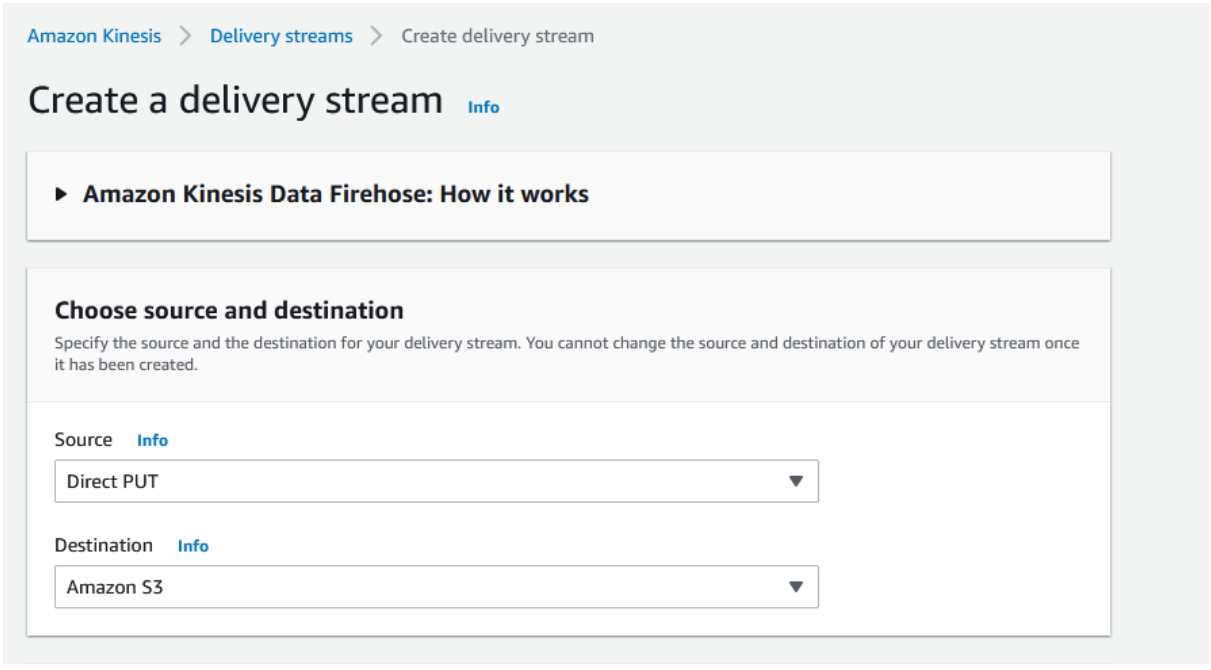

Selecting Destination as S3, create a new bucket. The bucket should have a unique name.

Change the buffer size and buffer interval as needed.

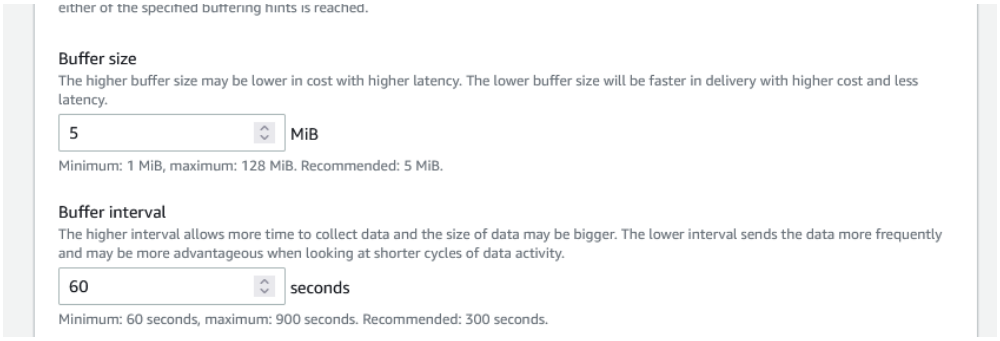

Create a new IAM role –

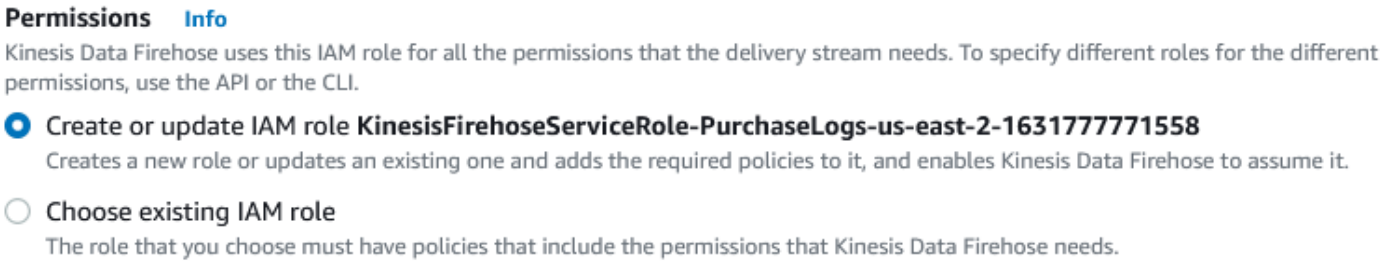

2. Now that the delivery stream is set up, we need to set up an EC2 instance to feed data into the stream.

Click on the EC2 service.

Then launch instance.

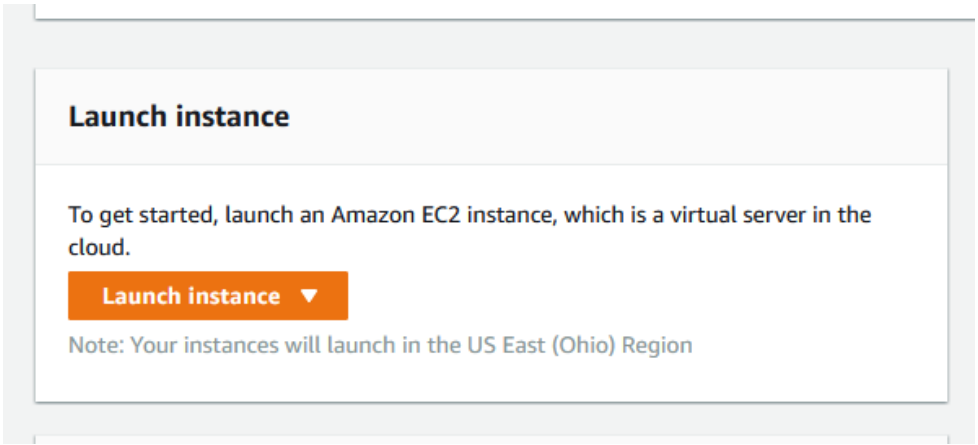

Select the Amazon Linux 2 AMI

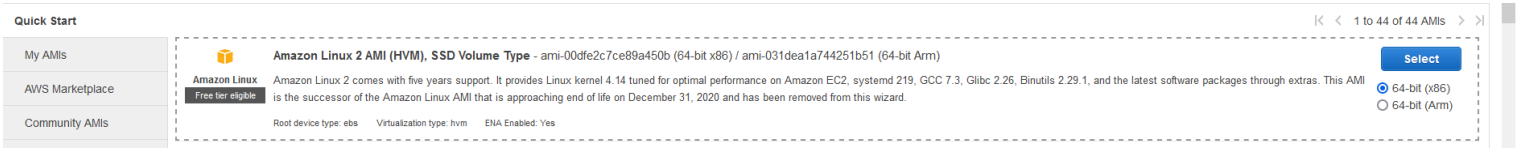

Select t2 micro free tier.

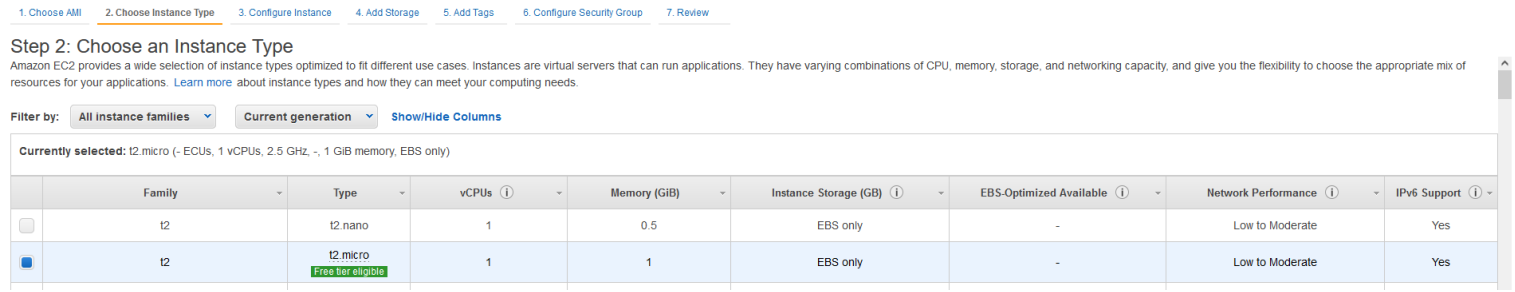

Click Launch again on the next screen. It would ask you to select an existing key pair if you haven't created one already.

Create one.

Finish creating the instance.

Connect to the instance form the screen after selecting it here –

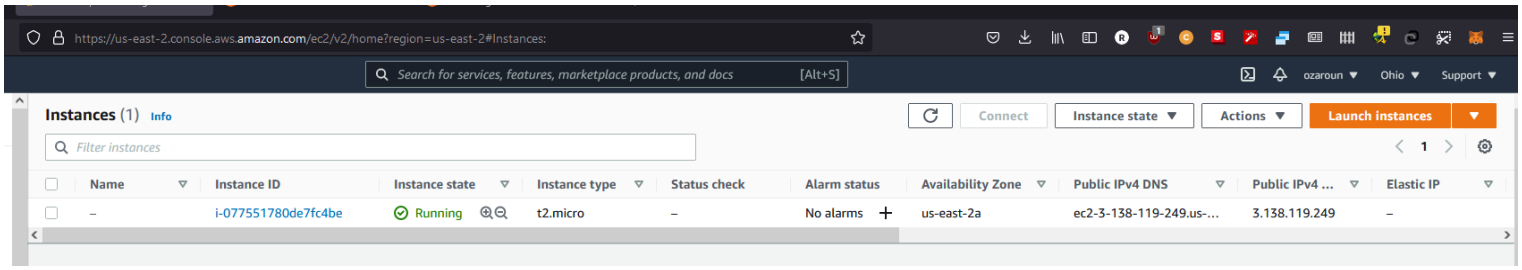

Connect to the EC2 using SSH

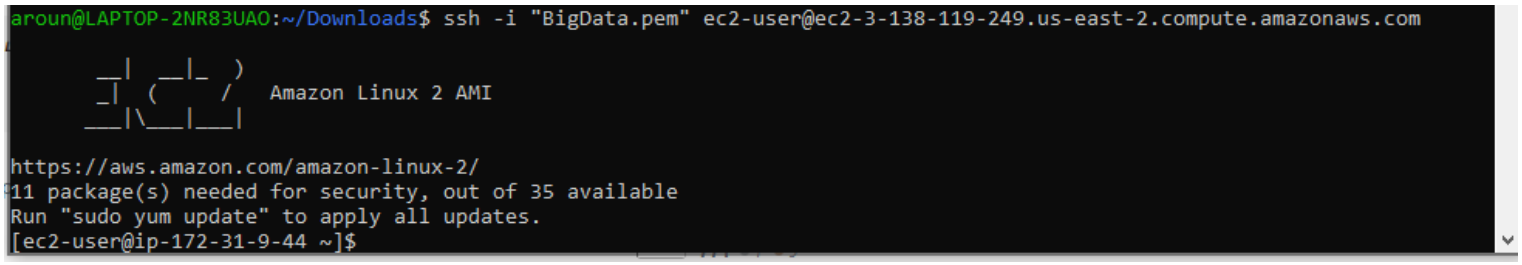

Install Kinesis agent with the command –

sudo yum install -y aws-kinesis-agent

Download the log generator for the purpose of the project – wget https://media.sundog-soft.com/AWSBigData/LogGenerator.zip

Unzip it – unzip LogGenerator.zip

Change permissions for the python file

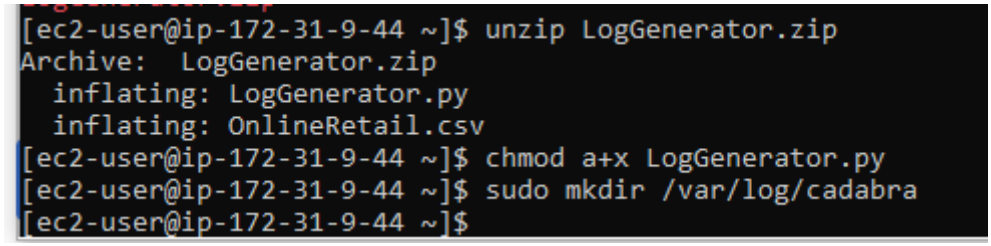

Change Kinesis Agent setting in -

```
GNU nano 2.9.8                                          agent.json

{
  "cloudwatch.emitMetrics": true,
  "kinesis.endpoint": "",
  "firehose.endpoint": "firehose.us-east-2.amazonaws.com",

  "flows": [
    {
      "filePattern": "/var/log/cadabra/*.log",
      "deliveryStream": "PurchaseLogs"
    }
  ]
}
```

In EC2 console screen select the instance, go to actions, go to security, go to IAM Role.

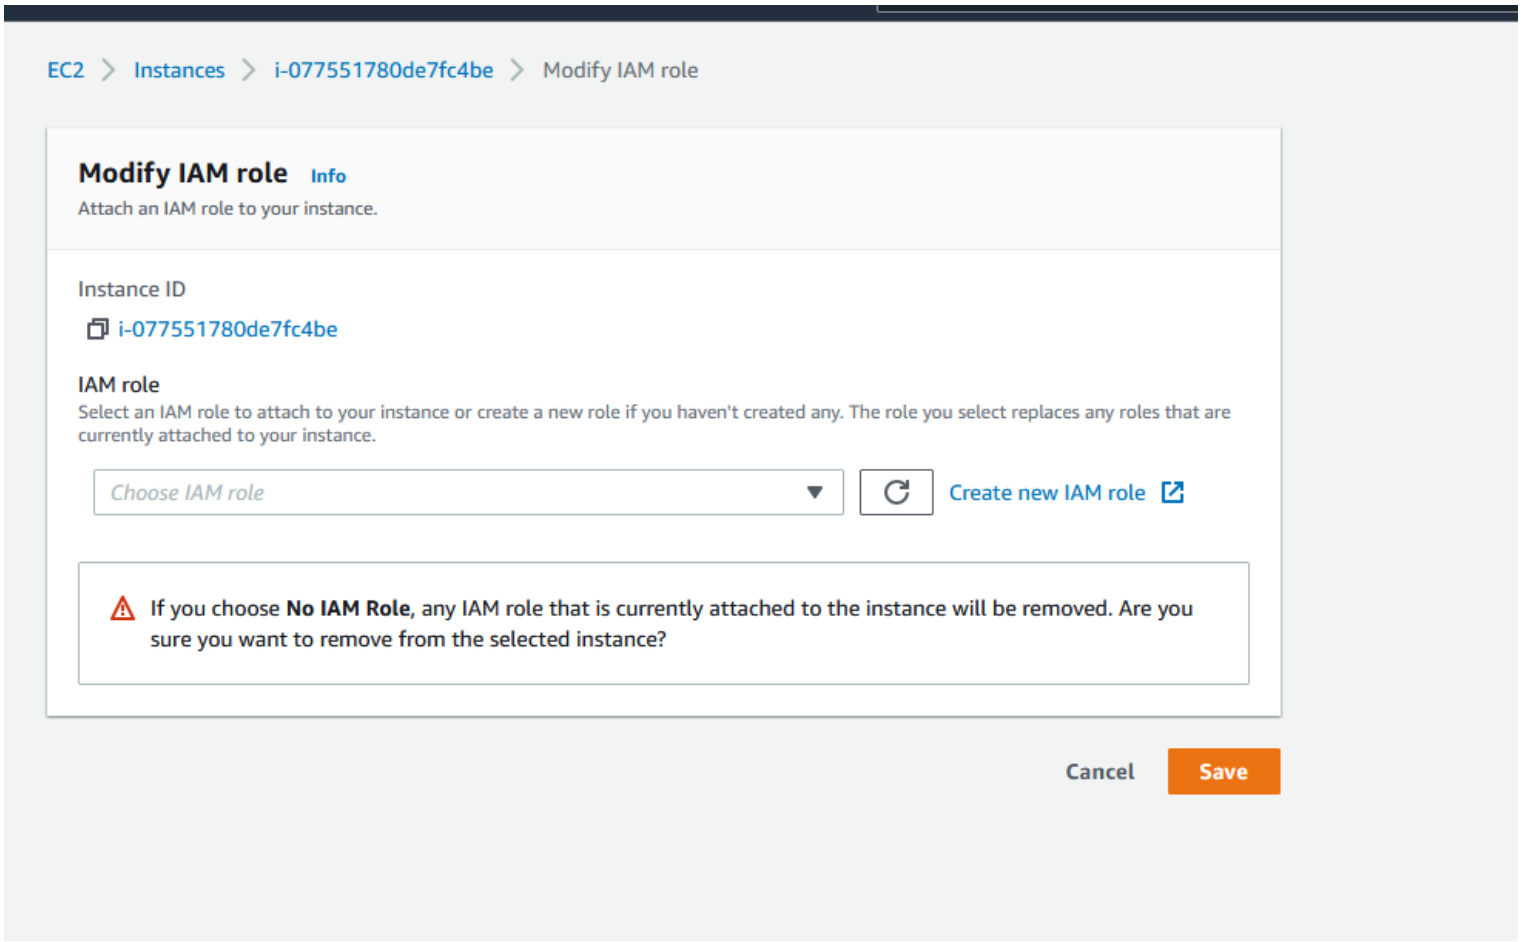

Click Create new IAM role.

Select Create Role

Select EC2. Select AdministratorAccess for ease. Name it and then create role.

Back to the EC2 screen where we selected the create new IAM role, use the refresh button and select the new IAM role and save.

Return to the Kinesis agent json file.

Start the kinesis agent –

```
[ec2-user@ip-172-31-9-44 aws-kinesis]$ sudo service aws-kinesis-agent start
Starting aws-kinesis-agent (via systemctl):                [  OK  ]
[ec2-user@ip-172-31-9-44 aws-kinesis]$
```

To make it start automatically

```
[ec2-user@ip-172-31-9-44 aws-kinesis]$ sudo service aws-kinesis-agent start
Starting aws-kinesis-agent (via systemctl):                [  OK  ]
[ec2-user@ip-172-31-9-44 aws-kinesis]$ sudo chkconfig aws-kinesis-agent on
```

Start the stream –

```
[ec2-user@ip-172-31-9-44 aws-kinesis]$ cd ~
[ec2-user@ip-172-31-9-44 ~]$ sudo ./LogGenerator.py 500000
Writing 500000 lines starting at line 0

Wrote 500000 lines.

[ec2-user@ip-172-31-9-44 ~]$
```

Check for the log and also can check the live number of logs generated using the tail function

```
[ec2-user@ip-172-31-9-44 aws-kinesis]$ cd ~
[ec2-user@ip-172-31-9-44 ~]$ sudo ./LogGenerator.py 500000
Writing 500000 lines starting at line 0

Wrote 500000 lines.

[ec2-user@ip-172-31-9-44 ~]$ cd /var/log/cadabra/
[ec2-user@ip-172-31-9-44 cadabra]$ ls
20210916-082658.log
[ec2-user@ip-172-31-9-44 cadabra]$ tail -f /var/log/aws-kinesis-agent/aws-kinesis-agent.log
2021-09-16 08:31:08.318+0000  (FileTailer[fh:PurchaseLogs:/var/log/cadabra/*.log].MetricsEmitter RUNNING) com.amazon.kin
esis.streaming.agent.tailing.FileTailer [INFO] FileTailer[fh:PurchaseLogs:/var/log/cadabra/*.log]: Tailer Progress: Tail
er has parsed 500000 records (42036722 bytes), transformed 0 records, skipped 0 records, and has successfully sent 50000
0 records to destination.
```

Data loaded into S3 in 5mb chunks

Name	Type	Last modified	Size	Storage class
PurchaseLogs-1-2021-09-16-08-27-01-bc9ab9e6-f5d4-4fbc-ac8d-e58506e27ddb	-	September 16, 2021, 01:27:06 (UTC-07:00)	4.7 MB	Standard
PurchaseLogs-1-2021-09-16-08-27-03-06abffcf-17ee-4044-9ca1-d69458f04f33	-	September 16, 2021, 01:27:18 (UTC-07:00)	4.7 MB	Standard
PurchaseLogs-1-2021-09-16-08-27-17-b89d2d8a-43fe-494b-a9bc-21d9c28b9ba3	-	September 16, 2021, 01:27:19 (UTC-07:00)	4.7 MB	Standard
PurchaseLogs-1-2021-09-16-08-27-18-12ab9e8f-a106-4c01-9144-7079eea4856f	-	September 16, 2021, 01:27:22 (UTC-07:00)	4.7 MB	Standard
PurchaseLogs-1-2021-09-16-08-27-20-9328e1d2-8ddb-43eb-a0ef-6584a55fcc0d	-	September 16, 2021, 01:27:34 (UTC-07:00)	4.7 MB	Standard
PurchaseLogs-1-2021-09-16-08-27-33-4a85aa09-02f7-4835-b8ee-7fb73d2e1654	-	September 16, 2021, 01:27:35 (UTC-07:00)	4.7 MB	Standard
PurchaseLogs-1-2021-09-16-08-27-33-b2fe553b-c388-4d72-b07b-bdff44789ef8	-	September 16, 2021, 01:27:59 (UTC-07:00)	4.7 MB	Standard
PurchaseLogs-1-2021-09-16-08-27-57-21b44288-b483-4d60-86db-ca31de1e0232	-	September 16, 2021, 01:27:59 (UTC-07:00)	4.7 MB	Standard
PurchaseLogs-1-2021-09-16-08-27-58-f8938433-47b6-418b-b990-b6ad7494ef49	-	September 16, 2021, 01:29:00 (UTC-07:00)	2.8 MB	Standard

So now we have a S3 data lake ready.

The end goal is the create an order history app. So that a user is able to access their order history on cadabra.com or whatever client they are using.

So we already have server logs from the EC2 instance.

Now instead of Firehose we're gonna publish it into Kinesis Data Stream so that the data is accessible in real time

3. Connecting the EC2 to a Kinesis Data Stream

On the Kinesis Console. Create Data Stream.

Name it and select the number of shards you are going to use. We are going to select 1.

Go to the ec2 instance and configure the kinesis agent json file for the logs to be sent to this stream.

```
GNU nano 2.9.8                                    agent.json

{
  "cloudwatch.emitMetrics": true,
  "kinesis.endpoint": "kinesis.us-east-2.amazonaws.com",
  "firehose.endpoint": "firehose.us-east-2.amazonaws.com",

  "flows": [
{
      "filePattern": "/var/log/cadabra/*.log",
      "kinesisStream": "CadabraOrders",
      "partitionKeyOption": "RANDOM",
      "dataProcessingOptions": [
         {
            "optionName": "CSVTOJSON",
            "customFieldNames": ["InvoiceNo", "StockCode", "Description", "Quantity", "InvoiceDate", "UnitPrice", "Customer", "Country"]
         }
      ]
    },
    {
      "filePattern": "/var/log/cadabra/*.log",
      "deliveryStream": "PurchaseLogs"
    }
  ]
}
```

Now the data is going to the S3 from the firehose and to the Kinesis Data Stream

4. Integrating DynamoDB

From the AWS console, select DynamoDB.

Create a table – CadabraOrders. Give Partition key – CustomerID of type Number. Add sort key – OrderID of type String.

On the EC2 instance, install the boto3 library – pip3 install boto3

Now we need to create some credentials files so that boto3 knows how to log on to S3.

```
[ec2-user@ip-172-31-9-44 ~]$ pwd
/home/ec2-user
[ec2-user@ip-172-31-9-44 ~]$ mkdir .aws
[ec2-user@ip-172-31-9-44 ~]$ cd .aws/
[ec2-user@ip-172-31-9-44 .aws]$ nano credentials
[ec2-user@ip-172-31-9-44 .aws]$
```

```
ec2-user@ip-172-31-9-44:~/.aws
  GNU nano 2.9.8

[default]
aws_access_key_id=[redacted]
aws_secret_access_key=[redacted]
```

Create config file in the same place.

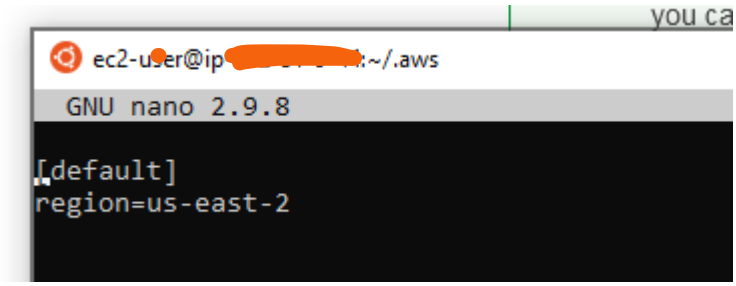

Now, downloading the consumer script in the home directory.

wget https://media.sundog-soft.com/AWSBigData/Consumer.py

Make the script executable - chmod a+x Consumer.py

From a second ssh into the EC2 instance, run LogGeneratory.py file for 100 records - ./LogGenerator.py 100

After a minute or so the first console where Consumer.py is running would pick up the records for the Kinesis Stream.

Now, we have created a system that works from end-to-end. We monitor new information being uploaded into the log directory of the EC2 host. The kinesis stream picks that data up. Another app in our case the Comsumer.py inserts that data into a DynamoDB table. We can now imagine a mobile application that talks directly to the DynamoDB instance and returns the information.

4. Integrating AWS Lambda

Right now we have a consumer script that is running on our EC2 instance playing the role of a lambda function. It is not a very scalable function.

Our Lambda function in terms of security would need to consume data from the kinesis stream and write data into the DynamoDB.

Let's create an IAM role for lambda.

Go to IAM console -> Go to Roles -> Click create role -> Select Lambda -> Attach the following permission to read from the Kinesis Stream and Write to the DynamoDB : AmazonKinesisReadOnlyAccess , AmazonDynamoDBFullAccess -> Click Next -> Skip tags -> Give the role a name (in our case CadabraOrders) -> Create Role

Now back to the AWS console. Select Lambda.

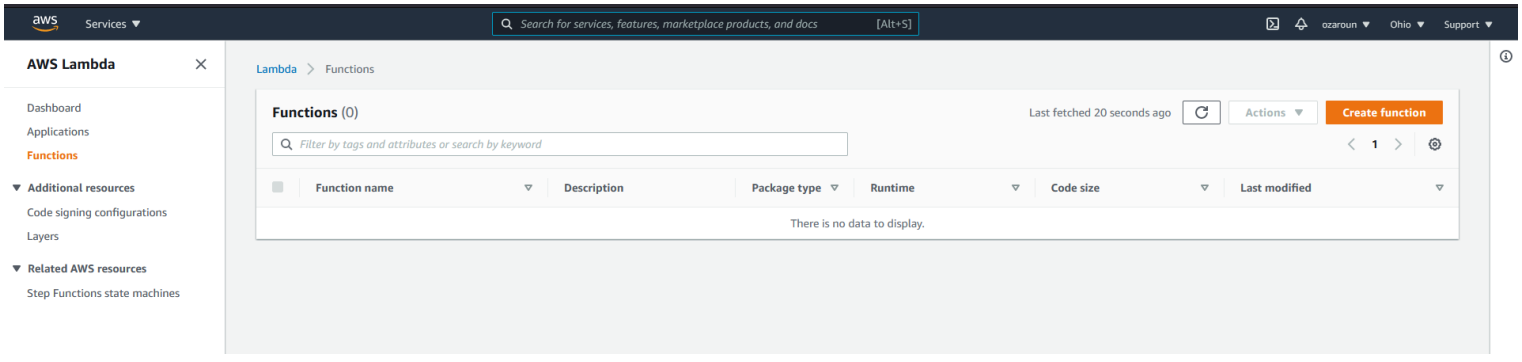

Select the create function button.

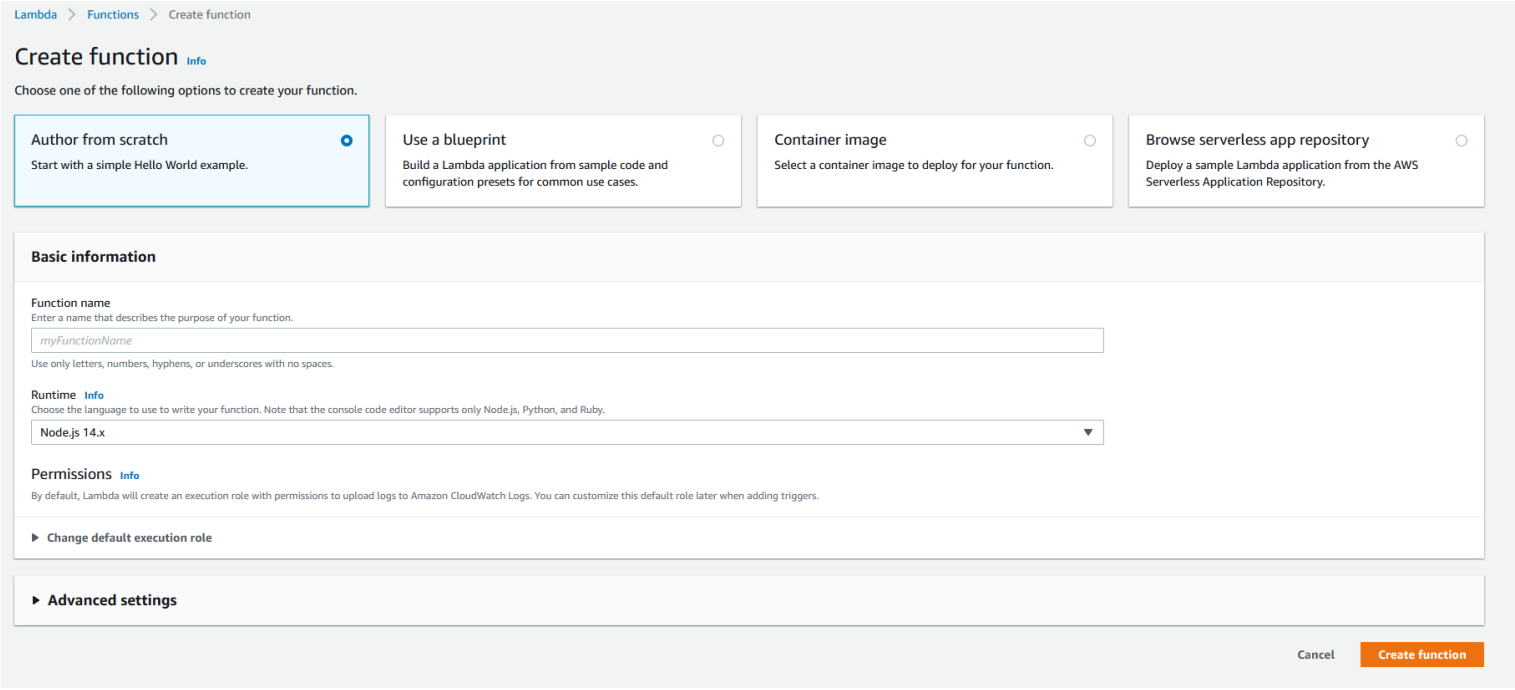

Give it a name (in our case ProcessOrders). Select the runtime environment – Python 3.x.

Select the existing IAM role that we create for it.

The Next Screen –

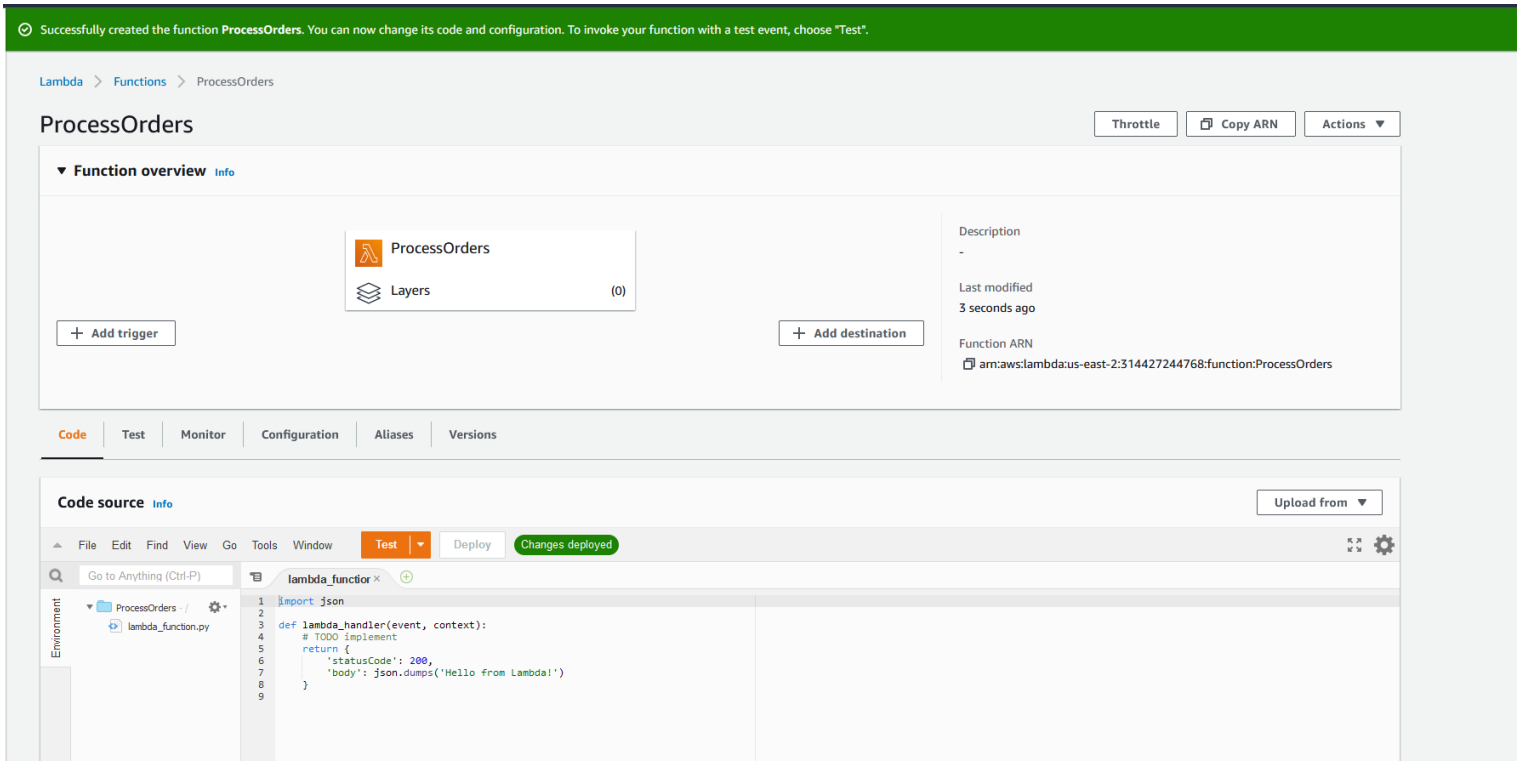

We want to add a trigger to feed data into this Lambda function from our Kinesis Stream.

Select Add Trigger -> Select Kinesis Streams -> Select the Kinesis Stream (in our case Cadabra Orders) -> Click Add.

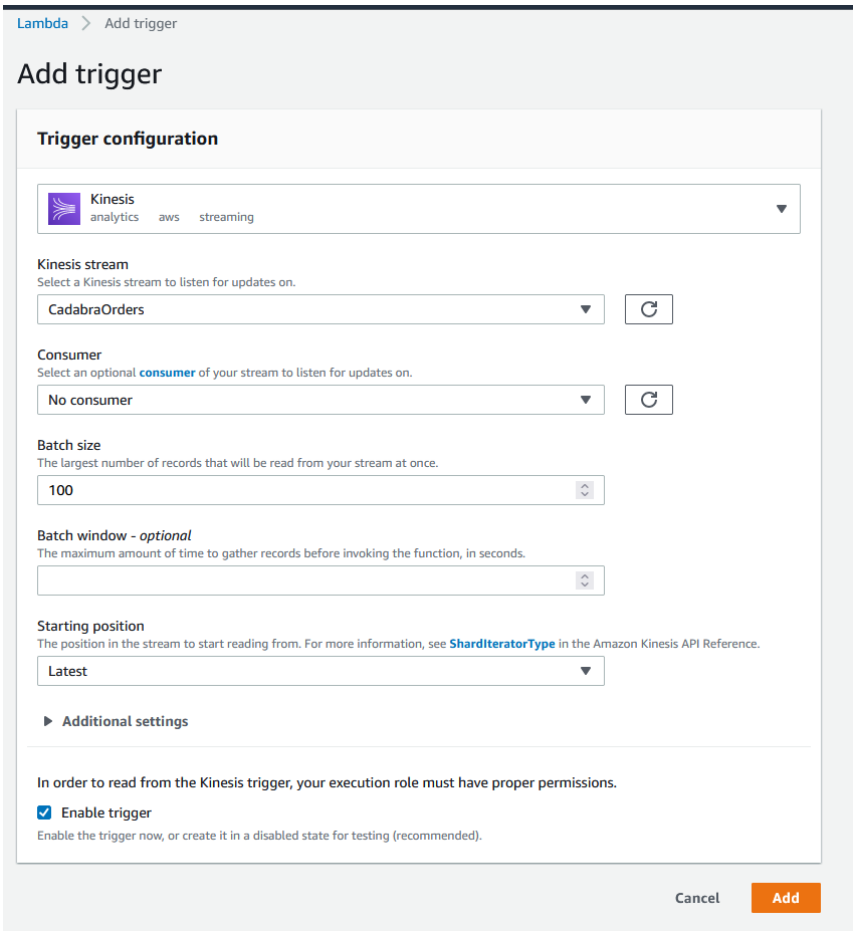

The Lambda function screen looks like –

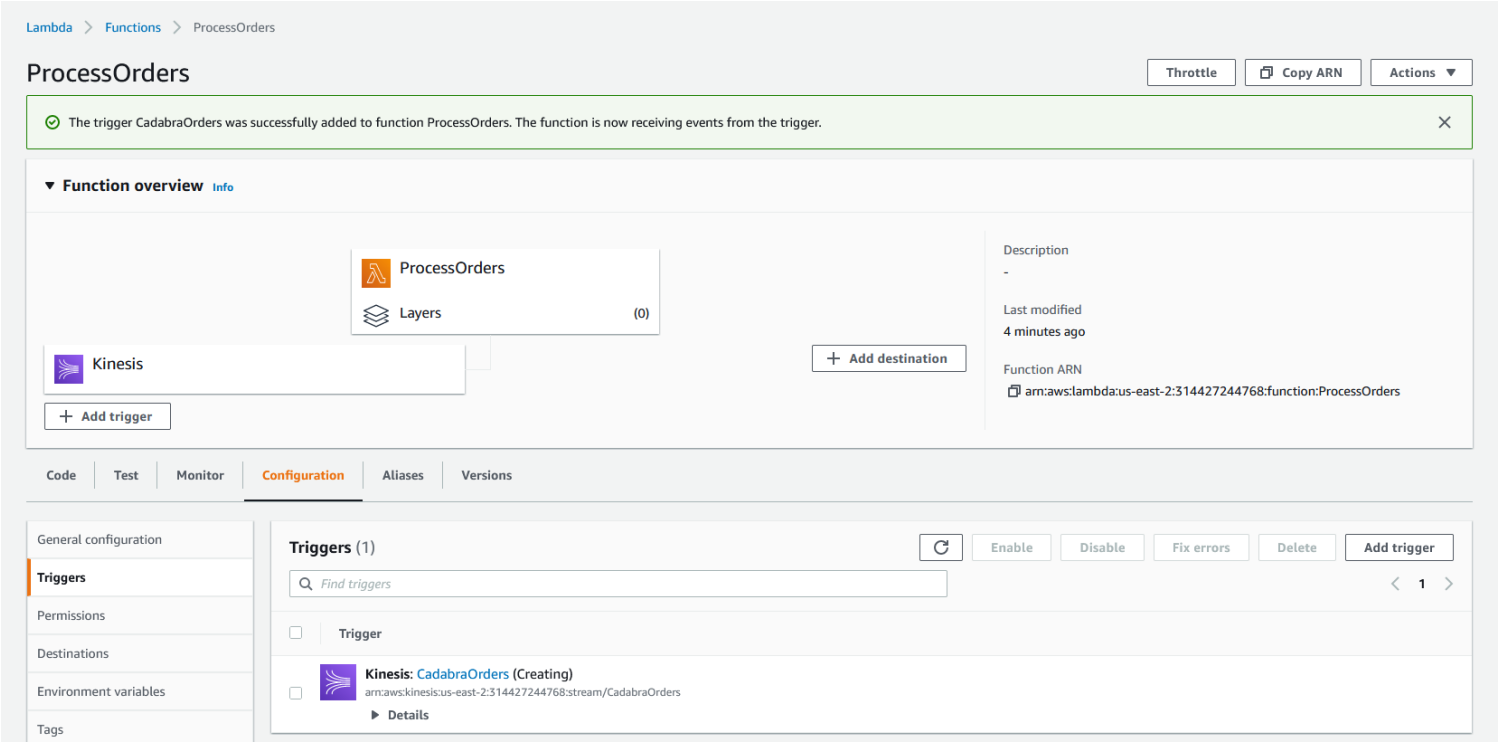

Click on code and add the following code –

```
import base64
import json
import boto3
import decimal
import uuid

def lambda_handler(event, context):
    item = None
    dynamo_db = boto3.resource('dynamodb')
    table = dynamo_db.Table('CadabraOrders')
    decoded_record_data = [base64.b64decode(record['kinesis']['data']) for record in event['Records']]
    deserialized_data = [json.loads(decoded_record) for decoded_record in decoded_record_data]

    with table.batch_writer() as batch_writer:
        for item in deserialized_data:
            # We've added a try / except block here to deal with invalid input rows more gracefully.
            # Be aware there are stretches of the input data that have no customer ID's at all,
            # keep trying the LogGenerator script to get past that if you run into it.
            try:
                invoice = item['InvoiceNo']
                customer = int(item['Customer'])
                orderDate = item['InvoiceDate']
                quantity = item['Quantity']
                description = item['Description']
                unitPrice = item['UnitPrice']
                country = item['Country'].rstrip()
                stockCode = item['StockCode']

                # Construct a unique sort key for this line item
                # We've added a uuid at the end as there is some duplicate invoice/stockcode
                # data in our sample data.
                orderID = invoice + "-" + stockCode + "-" + uuid.uuid4().hex

                batch_writer.put_item(Item = {
                        'CustomerID': decimal.Decimal(customer),
                        'OrderID': orderID,
                        'OrderDate': orderDate,
                        'Quantity': decimal.Decimal(quantity),
                        'UnitPrice': decimal.Decimal(unitPrice),
                        'Description': description,
                        'Country': country
                    }
                )
                print("Wrote item into batch.")
            except:
                print("Error processing invalid input row.")
```

Save the changes. Deploy the code.

We don’t explicitly have to mention DynamoDB in the destination, our script is doing that for us.

Lets ssh into our EC2 instance.

Run the command - sudo ./LogGenerator.py 100

Restart the aws-kinesis-agent - sudo service aws-kinesis-agent restart

The newer records for some reason didn't show up on the DynamoDB as quickly as I had expected but they did show up after a few minutes.

5. EMR MapReduce –

We'll build a product recommendations system for "kadabra.com" , the application we have been building. We already have deployed a firehose that dumps data from the EC2 instance to the S3 bucket. We can now deploy an EMR Cluster and using Apache Spark and MLlib we can generate recommendations based on order data in S3.

Head over to the AWS console and click on EMR -> Click on create cluster ->

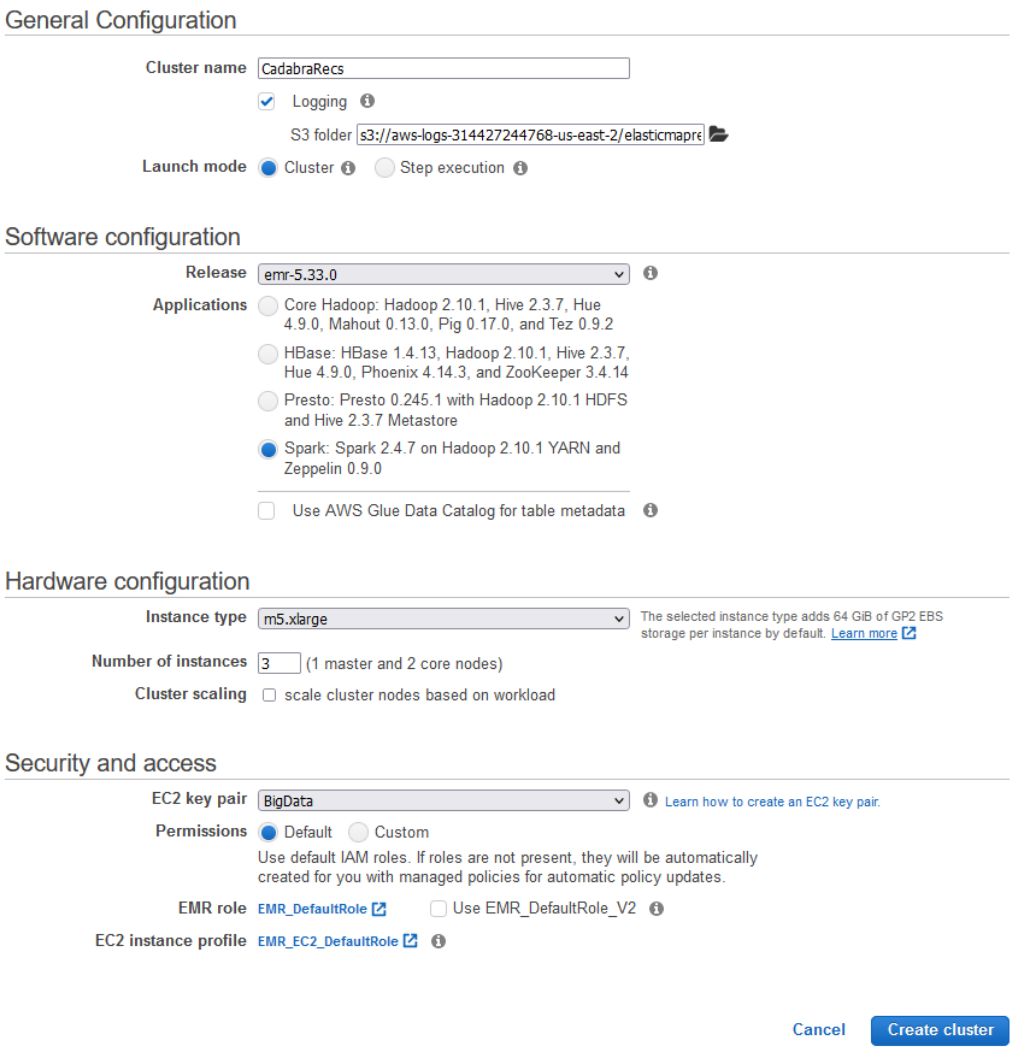

The EMR cluster does not come under the free tier. Create the cluster.